

DAT200 CA5 2022

Kaggle username: Mohamed Atteyeh

Imports

```
In [159... import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import re
import seaborn as sns
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, PolynomialFeatures, MinMaxScaler
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.linear_model import LinearRegression, Ridge, ElasticNet, Lasso, Perceptron
from sklearn.feature_selection import SelectKBest, f_regression
from sklearn.svm import SVC, LinearSVC
from sklearn.decomposition import PCA
from sklearn.model_selection import StratifiedKFold, cross_val_score
from sklearn.model_selection import GridSearchCV
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.svm import LinearSVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import TomekLinks
```

Reading data

```
In [160... Tr_data = pd.read_pickle ('train.pkl')


# Cleaning the column names
columns = Tr_data.columns
for col in columns:
    Tr_data = Tr_data.rename(
        columns={
            # strip out parentheses, and their contents
            col: re.sub(r'\(.*', '', col)
                .strip()
                .replace(' ', '_')
                .replace('-', '_')
                .lower() # lowercase the column name
        }
    )
df = Tr_data.copy()

# Removal of all values that has missing in to be able to see the visualisation
missing_values = df[df.eq('missing').any(1)]
df = df.drop(df.index[list(missing_values.index)])
```


Data exploration and visualisation

```
In [161... cat_columns = ['year', 'month', 'hour', 'holiday', 'weekday', 'working_day',
                    'temperature', 'feels_like_temperature',
                    'humidity', 'windspeed', 'rental_bikes_count']
for i in cat_columns:
    df[i] = df[i].astype(float)


corr = df[list(cat_columns)].corr()
plt.figure(figsize = (20,10))
sns.heatmap(corr,annot = True, cmap = 'BrBG')
plt.show()
```



```
In [162... plt.figure(figsize = (20,10))
sns.scatterplot(x = df['hour'],y = df['rental_bikes_count'],hue = df['season'])
plt.show()
```



```
In [163... plt.figure(figsize = (20,10))
sns.scatterplot(x = df['temperature'],y = df['rental_bikes_count'],hue = df['season'])
plt.show()
```



Data cleaning

```
In [164... # Choose label encoding, since there are multiple columns with the same build
df = df.replace(to_replace=['Heavy rain, heavy snow or thunderstorm', 'Light
df = df.replace(to_replace=['Summer', 'Winter', 'Fall', 'Spring'], value = [0,1
```

Data preprocessing

```
In [165... X,y = df.drop('rental_bikes_count', axis = 1).copy(), df['rental_bikes_count
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.4,random_st

# PIPELINES
DT_regressor = make_pipeline(DecisionTreeRegressor(max_depth=10))
RF_regressor = make_pipeline(RandomForestRegressor(n_estimators=100, random_

# Testign and fitting the split
DT_regressor.fit(X_train,y_train)
RF_regressor.fit(X_train,y_train)

print('Train accuracy DT: {0:.5f}'.format(DT_regressor.score(X_train, y_trai
print('Test accuracy DT: {0:.5f}'.format(DT_regressor.score(X_test, y_test))

print('Train accuracy RF: {0:.5f}'.format(RF_regressor.score(X_train, y_trai
print('Test accuracy RF: {0:.5f}'.format(RF_regressor.score(X_test, y_test))

Train accuracy DT: 0.93511
Test accuracy DT: 0.88462
Train accuracy RF: 0.99134
Test accuracy RF: 0.93779
```

Bins Model

```
In [167... # Bins Model
dfc = df.copy()
dfc = pd.get_dummies(dfc, columns = ['hour', 'month', 'weekday'])
forest = RandomForestClassifier(criterion='gini',n_estimators=60,random_stat
svm = make_pipeline(SVC(kernel='rbf', C=20, random_state=100))

binned = []
for i in range(2,11):
    y_bin= pd.cut(dfc['rental_bikes_count'], bins=i, labels=False)
```

```

Xc = dfc.drop('rental_bikes_count', axis = 1).copy()
Xc_train,Xc_test,yc_train,yc_test = train_test_split(Xc,y_bin,test_size=

forest.fit(Xc_train, yc_train)
FTrain = forest.score(Xc_train, yc_train)
FTest = forest.score(Xc_test, yc_test)

sc = StandardScaler()
sc.fit(Xc_train)

Xc_train_sc = sc.transform(Xc_train)
Xc_test_sc = sc.transform(Xc_test)

svm.fit(Xc_train_sc, yc_train)
svmtrain = svm.score(Xc_train_sc, yc_train)
svmtest= svm.score(Xc_test_sc, yc_test)

binned.append ([i,svmtrain,svmtest,FTrain,FTest])

table = pd.DataFrame(binned , columns=['bins number','SVM Train','SVM Test',
table

```

Out[167]:

	bins number	SVM Train	SVM Test	Random Forest Train	Random Forest Test
0	2	0.995464	0.963861	0.999433	0.960459
1	3	0.993622	0.927296	0.992771	0.911139
2	4	0.989086	0.890944	0.986676	0.875850
3	5	0.986393	0.844600	0.981999	0.829294
4	6	0.981715	0.811862	0.975478	0.783163
5	7	0.982282	0.780400	0.981999	0.754464
6	8	0.978313	0.758929	0.973919	0.740009
7	9	0.977746	0.731080	0.972927	0.701956
8	10	0.973777	0.702168	0.975195	0.677934

Kaggle submission

```

In [ ]: RF_regressor = make_pipeline(RandomForestRegressor(n_estimators=100, random_

kgl_X,kgl_y = df.drop('rental_bikes_count', axis = 1).copy(), df['rental_bik

RF_regressor.fit(kgl_X,kgl_y)
kgl_pred = RF_regressor.predict(df)

output = pd.DataFrame({'idx': df.index,'Rental bikes count': kgl_pred})
output.to_csv('CA5_submission_RF', index = False)

# MY Results were 93.134 using the Random Forest Regressor

```